

RootJS: Node.js Bindings for ROOT 6

Theo Beffart, Maximilian Früh, Christoph Haas, Sachin Rajgopal,
Jonas Schwabe, Christoph Wolff, Marek Szuba

Steinbuch Centre for Computing, Karlsruhe Institute of Technology,
Hermann-von-Helmholtz-Platz 1, 76344 Eggenstein-Leopoldshafen, Germany

E-mail: Marek.Szuba@cern.ch

Abstract. We present rootJS, an interface making it possible to seamlessly integrate ROOT 6 into applications written for Node.js, the JavaScript runtime platform increasingly commonly used to create high-performance Web applications. ROOT features can be called both directly from Node.js code and by JIT-compiling C++ macros. All rootJS methods are invoked asynchronously and support callback functions, allowing non-blocking operation of Node.js applications using them. Last but not least, our bindings have been designed to platform-independent and should therefore work on all systems supporting both ROOT 6 and Node.js.

Thanks to rootJS it is now possible to create ROOT-aware Web applications taking full advantage of the high performance and extensive capabilities of Node.js. Examples include platforms for the quality assurance of acquired, reconstructed or simulated data, book-keeping and e-log systems, and even Web browser-based data visualisation and analysis.

1. Introduction

1.1. Background

The scientific software framework ROOT is a *de facto* standard tool in nuclear and particle physics [1]. It provides components for a wide range of purposes such as statistical data analysis, advanced visualisation, machine learning, Monte-Carlo simulations, persistence of data, parallel processing, and many more. Although it has been written mainly in C++ and typically interfaced with using macros (scripts) written in that language, ROOT bindings exist for other languages such as Python, R, or Ruby.

Among its numerous capabilities ROOT features a set of building blocks for graphical user interfaces. ROOT-based GUI applications include event viewers, control and monitoring applications, and data quality-assurance tools. However, such tools have traditionally been built as monolithic applications which require ROOT and all of its dependencies installed locally and, unless carefully designed for that purpose, offered minimal scalability.

A well-known alternative to monolithic tools are Web applications — client-server software in which the client runs in a Web browser. Unlike traditional software, Web applications naturally follow a distributed model of multiple tiers, each of them responsible for a different part of application logic. This model greatly improves scalability because different tiers can be extended independently. Moreover not only is the client component of a Web application typically lightweight, it also only requires the user to run a modern, standard-compliant Web browser — doing away with the need for local installation or creating builds for many different operating systems and hardware architectures. At the same time, modern Web standards and techniques greatly reduce constraints imposed by having to run an application inside a Web

browser. To name just three: Ajax allows for the the user interface to be dynamically updated inside the browser without the need for reloading the entire page, WebGL leverages local GPU power to render 3D graphics, and IndexedDB enables the use transactional local databases in applications.

1.2. Node.js and the MEAN Stack

One of the defining characteristics of a Web application is the *software stack* it is based on. Possibly the best known among these is LAMP, which originally consisted of the Linux operating system, the Apache HTTP Server, the MySQL relational database management system and the PHP programming language, and which along with its derivatives remains arguably the most popular model on the Internet [2].

An alternative, more recent Web-application stack is known as MEAN and consists of the NoSQL database MongoDB, the Web-application framework Express running on top of the server-side platform Node.js, and the front-end framework AngularJS [3]. It offers several advantages to both users and developers. For example: its components have been specifically designed for high performance, it naturally allows the client to perform more tasks than the mere rendering of content fetched from the server, and the use of a single programming language throughout the stack (JavaScript) simplifies both development and debugging.

Of the four MEAN components, the one directly relevant to the present paper is Node.js. It is an open-source and cross-platform JavaScript runtime environment based on the V8 JavaScript engine created by Google. The core runtime is minimalistic by design but can be extended with modules (for instance one to provide a HTTP server, which is how one most commonly uses Node in the server-side tier of Web applications), the management of which is simplified by the Node Package Manager (**npm**) and the corresponding public module directory [4]. An important feature of Node.js is its high performance, which it owes to the fact that it is almost entirely asynchronous — instead of blocking on I/O operations, it handles their results via callback functions.

1.3. Motivation

The motivation behind work described in the present paper has been to develop Open Source JavaScript bindings for ROOT, which we have called *rootJS*, which would allow it to be used in Node.js scripts and MEAN stack-based Web applications. The following sections of the manuscript discuss the architecture of these bindings, outlines some of the details of their implementation along with associated challenges, and discusses potential use cases. Finally, we discuss related work.

2. Overview of rootJS

2.1. Functional Requirements

The most important requirement we defined for our bindings was that they must be complete. Given the ROOT class library already numbers into thousands, is continuously extended by upstream developers, and can be supplemented at both compile and run time by third-party components such as those created by specific experiments, this requires the lists of classes and methods available through rootJS to be handled dynamically. On a related note, we also wanted to allow the users to load additional ROOT libraries at run time.

Secondly, the bindings must take the form of a Node.js module *i.e.* require no modifications to Node.js source code. Similarly they should communicate with ROOT using its standard APIs, again not requiring modifications of upstream code.

Next, the bindings should keep with the spirit of Node.js and provide asynchronous wrappers for at least the most common I/O functions. Ideally they should support asynchronous execution of all ROOT functions.

Finally, the bindings should allow the user to feed C++ code to the ROOT just-in-time-compiler to support re-use of existing code blocks. Needless to say rootJS must accurately reflect the state of ROOT following execution of such code, which provided another argument in favour of the dynamic approach.

2.2. Software Requirements

We have decided rootJS would only support ROOT 6 because its Low Level Virtual Machine (LLVM)-based C++ interpreter Cling offers many advantages, *e.g.* the aforementioned just-in-time compilation, over the one available in older ROOT versions.

Regarding operating-system compatibility, our goal has been for the bindings to support GNU/Linux on amd64 architecture. That said, completed bindings have been shown to run correctly under Mac OS X and should in theory be compatible with all systems supported by both ROOT 6 and Node.js.

2.3. Challenges

Developing JavaScript bindings for a C++ library like ROOT is not a trivial task because of fundamental differences between the two languages. Most notably:

- JavaScript is a functional language with first-class functions; the latter feature is widely used for callbacks. In contrast, C++ is an imperative and object-oriented language in which functions are not first-class citizens;
- they employ different type systems: C++ is strongly typed, JavaScript offers only limited dynamic type checking. One important consequence of the above is that a JavaScript engine will not overload functions depending on the type of its arguments, which is widely applied in ROOT (see *e.g.* `TH1::Fill(Double_t, Double_t)` *vs* `TH1::Fill(const char*, Double_t)`);
- on a related note, the two languages support different data types. For instance: JavaScript only has a single numeric type; only C++ supports enumerated types; strings are a primary data type in JavaScript but objects in C++;
- C++ supports low-level memory access and can require the user to implement their own memory management. JavaScript has neither;
- unlike C++, JavaScript objects are classless.

Fortunately both Node.js and ROOT provide ways in which these differences can be reconciled in order to create an *adapter* between the two:

- the C++ API of the V8 engine allows exposure of C++ objects to Node.js, can map JavaScript objects to C++ classes, and provides callback handling;
- ROOT features extensive run-time introspection capabilities, covering both individual functions/objects as well as general information about classes, name spaces, and global and member variables.

3. Architecture

The following elements have been implemented in rootJS in order for it to meet our requirements.

To begin with, the bindings recursively seek and expose ROOT classes and name spaces. This is done at initialisation time, after the loading of a new library, after execution of JIT-compile code, or upon user request. The name-space hierarchy of ROOT is preserved and reproduced by a tree of properties attached to the top-level object of the bindings.

Secondly, existing objects, functions and variables must be wrapped in appropriate *proxy objects* which are then exposed by V8 to JavaScript; similar encapsulation is also established for

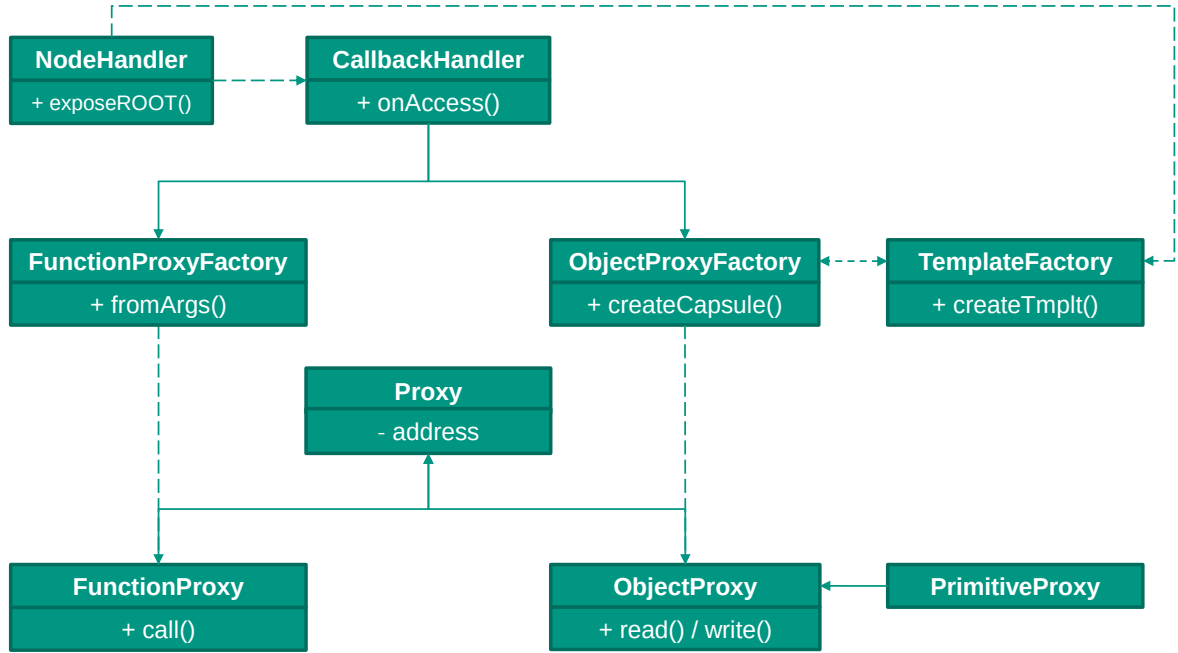


Figure 1. Core architecture of rootJS.

newly created ROOT objects after the forwarding of respective constructor calls. These proxy objects are produced by respective *factories*, which also cache their products so that duplication of work is avoided *e.g.* when multiple instances of the same ROOT object are created. The proxies also normalise memory addresses in order to properly handle the use of pointers common in ROOT.

It is worth emphasising at this point that all read and write operations on exposed objects and variables take place directly in ROOT memory space. That way everything stays in sync all the time.

Furthermore, function calls must be provided asynchronous call context in order to avoid blocking. We have implemented concurrent execution of functions using *libuv*, an asynchronous-I/O library designed for Node.js for this very purpose [5]; *libuv* maintains a pool of worker threads and communicates with them using messages. The original idea to employ the threading API provided by ROOT (TThread) had to be abandoned due to limitations of the V8 engine, which requires all interaction with Node.js to take place in the main thread.

Finally, a handler exists which maintains associations between function calls and their respective callbacks.

A diagram of the core architecture of rootJS can be found in Figure 1.

4. Getting and Using rootJS

rootJS has been registered in the npm module directory so for most users who have already got Node.js and its package manager installed, getting rootJS grinds down to a single command:

```
npm install rootjs
```

Alternatively one can manually download the source code from our GitHub repository or its CERN GitLab mirror [6, 7].

Please note that rootJS is a native C++ module that has to be compiled before use, which requires a working C++ compiler as well as libuv header files. Naturally one must also have ROOT 6 installed.

After installation the module is loaded the standard Node.js way:

```
var root = require('rootjs');
```

after which all ROOT variables, name spaces and classes become available through the object `root`.

rootJS is ultimately a thin wrapper for ROOT itself so most usage examples would be the same (including operations producing graphics on the screen — at least under Linux, using rootJS does not prevent ROOT from communicating with the local X server), making it unnecessary to repeat them here. An exception to this are callbacks, which can always be passed to functions as their last argument. For example:

```
root.TFile.Open("foo.root", function (fin) {  
    fin.ls();  
});
```

will return immediately, then only invoke `ls()` on the file once it has been opened. One can of course still invoke functions without callbacks in which case they are executed the standard ROOT way, only returning once the call has been finished.

Finally, the loading of additional libraries works slightly differently so that rootJS can immediately update its lists of classes and variables. Instead of using `gSystem->Load()`, one does the following:

```
root.loadlibrary("libMathCore.so"); // load it...  
root.ROOT.Math.Pi(); // ...and use it! Note the preservation of name spaces
```

For additional information, please see our GitHub site [6].

5. Use Cases

Like ROOT itself, rootJS is a set of building blocks to create various applications rather than an actual application. Its direct purpose is simply to bring the capabilities of ROOT to another programming language (JavaScript), with emphasis on a particular domain (high-performance Web applications). That said, we would like to discuss at this point two possible use cases for rootJS which have inspired us to initiate this project.

On one hand, rootJS could be used to implement a Web-based event viewer for a HEP experiment. Such viewers are commonly used for the monitoring of data taking by experiments as well as basic quality assurance, and they are not infrequently implemented using ROOT. There are, however, two disadvantages of using an event viewer in the form of a standalone application:

- limited portability — it must be installed on the target machine, along with ROOT and all other dependencies. This is at present an issue especially for applications requiring ROOT 6 due to its still-limited availability (*e.g.* lack of native Windows versions);
- such a viewer is usually located close to the data. This is particularly important for live monitoring (in which case the source of data is typically isolated from the Internet) but can also be an issue in off-line mode (in which case the data might be too big to move around quickly).

Using a modern, multi-tiered Web application instead can improve the situation: on one hand all the end user needs is a Web browser, on the other they can be located essentially anywhere because only the back-end tier has to be located near the data source. There has been growing interest in Web-based event browsers in the particle-physics community (see *e.g.* [8] in these proceedings) — and rootJS enables direct integration of ROOT into one of the increasingly popular back-end platforms.

On the other hand, rootJS makes it possible to employ the capabilities of ROOT for virtually any sort of analytics in MEAN stack-based Web applications. It could be used to integrate machine learning, statistical analysis, linear algebra ... but also access to XRootD, PROOF and so on. In our particular case, we plan to take advantage of ROOT in Web-based analysis of data from Earth-observing climatology satellites [9].

6. Related Work

In this section we would like to briefly introduce three examples of related work — JavaScript ROOT, ROOT `THttpServer`, and the ROOT kernel for Jupyter Notebook — and compare them to rootJS. These examples have been chosen not only because they are relevant to the use cases described above but also because they are now part of the core ROOT distribution.

6.1. JavaScript ROOT

In the words of its authors, “JavaScript ROOT provides interactive ROOT-like graphics in the web browsers. Data can be read and displayed from binary and JSON ROOT files. JSROOT implements user interface for `THttpServer` class.” It is a JavaScript reimplementaion of the aforementioned components of ROOT.

This description alone makes it clear that JSROOT and rootJS are complementary rather than competing solutions — JSROOT is a *front-end* component of a Web application whereas our bindings would be used on the *back-end* tier. It would be entirely feasible to create an application in which users interact with the former but in which all the under-the-bonnet processing is handled through the latter.

6.2. ROOT `THttpServer`

As mentioned above, ROOT 6 contains an integrated HTTP server. `THttpServer` is based on the embedded C/C++ server `CivetWeb` and continues to be extended with additional features [10, 11]. A question might therefore arise: why bother with Node.js and rootJS if the same can be achieved with ROOT itself?

Leaving aside the subjects of performance (which would have to be carefully benchmarked before any conclusions could be made) or individual preferences, we believe the key features in favour of rootJS are scalability and the wide range of modules available for Node.js which extend its capabilities as a Web server. `CivetWeb`, and by extension `THttpServer`, is a capable yet lightweight tool which for a lot of applications will be more than enough. However, what if you wanted to migrate your back-end server to HTTP2? Enable U2F-compliant two-factor authentication? Move from RESTful HTTP to WebSockets? Or perhaps just replace your single Web server with a high-availability cluster? In case of Node.js, all of these features are already supported — in some cases as easily as by a drop-in replacement of a module.

6.3. Jupyter Notebook

Jupyter Notebook is a Web application designed for creation and sharing of documents containing live code (in over 40 languages), equations, visualizations and explanatory text [12]. A ROOT kernel has been developed for Jupyter which enhances the latter with support for C++ ROOT code [13]. Moreover, by using the Python 2 kernel and the ROOT Python bindings (PyROOT) one interface ROOT with the numerous data-science Python modules such as SciPy.

While there is a certain overlap between Jupyter Notebook and rootJS, the two are considerably different: the former is a complete Web application featuring a sophisticated interactive interface, whereas the latter is a library interface aimed at scripted and batch processing on the back-end tier. Ultimately, they cover radically different use cases.

It is also worth noting at this point that Project Jupyter also features a JavaScript kernel that internally uses Node.js. Using rootJS with this kernel one can therefore grant ROOT notebooks access to the Node.js module ecosystem.

7. Conclusions

We have developed cross-platform Node.js bindings for ROOT 6, called rootJS, which makes it possible to integrate capabilities of ROOT into high-performance Web applications based on the MEAN stack. The bindings provide access to all classes and variables of ROOT, including those provided by additional libraries loaded at run time. They also allow the user to invoke all ROOT functions asynchronously with callbacks.

Acknowledgements

We would like to express our thanks to the authors of PyROOT, which has proven an extremely helpful reference for the Cling API.

The bulk of this work was done as part of the Software Engineering Practice class for Computer Science students of Karlsruhe Institute of Technology.

This work was conducted within the scope of the project “Large-Scale Data Management and Analysis”, funded by the German Helmholtz Association [14].

References

- [1] Brun R and Rademakers F 1997 *Nucl. Instrum. Meth.* **A389** 81–86
- [2] Kunze M 1998 *c't Archiv* 230 (in German)
- [3] Karpov V 2013 The MEAN Stack: MongoDB, ExpressJS, AngularJS and Node.js (last visited: 2017-02-20) URL <https://www.mongodb.com/blog/post/the-mean-stack-mongodb-expressjs-angularjs-and>
- [4] node package manager (last visited: 2017-02-20) URL <https://www.npmjs.com/>
- [5] libuv | Cross-platform asynchronous I/O (last visited: 2017-02-20) URL <http://libuv.org/>
- [6] URL <https://github.com/rootjs/rootjs>
- [7] URL <https://gitlab.cern.ch/rootjs/rootjs>
- [8] Mc Cauley T A browser-based event display for the CMS Experiment at the LHC using WebGL *these proceedings*
- [9] Szuba M, Ameri P, Grabowski U, Meyer J and Streit A 2016 A distributed system for storing and processing data from earth-observing satellites: System design and performance evaluation of the visualisation tool *IEEE/ACM 16th International Symposium on Cluster, Cloud and Grid Computing, CCGrid 2016, Cartagena, Colombia, May 16-19, 2016* (IEEE Computer Society) pp 169–174 ISBN 978-1-5090-2453-7 URL <http://dx.doi.org/10.1109/CCGrid.2016.19>
- [10] CivetWeb: Embedded C/C++ web server (last visited: 2017-02-20) URL <https://github.com/civetweb/civetweb>
- [11] Naumann A Status and Evolution of ROOT *these proceedings*
- [12] Kluyver T, Ragan-Kelley B, Pérez F, Granger B E, Bussonnier M, Frederic J, Kelley K, Hamrick J B, Grout J, Corlay S, Ivanov P, Avila D, Abdalla S, Willing C and et al 2016 Jupyter notebooks - a publishing format for reproducible computational workflows *Positioning and Power in Academic Publishing: Players, Agents and Agendas, 20th International Conference on Electronic Publishing, Göttingen, Germany, June 7-9, 2016*. ed Loizides F and Schmidt B (IOS Press) pp 87–90 ISBN 978-1-61499-648-4 URL <http://dx.doi.org/10.3233/978-1-61499-649-1-87>
- [13] Piparo D, Tejedor Saavaedra E and Mato Vila P The New ROOT Interface: Jupyter Notebooks *these proceedings*
- [14] Jung C, Gasthuber M, Giesler A, Hardt M, Meyer J, Rigoll F, Schwarz K, Stotzka R and Streit A 2014 *J. Phys. Conf. Ser.* **513** 032047